

Large Solution Management – Everything Will Be Fine, Right?

The February 2026 talk that started the situational-picture series

Version 1.3 · As of 2026-08-15 · Talk: SAFe Treff, Feb 26, 2026, Frankfurt am Main · Precedes all memoranda of the series

About this document: A curated write-up of the talk “**Large Solution Management – Everything Will Be Fine, Right?**”, delivered by **Robert Seebauer on February 26, 2026, at the SAFe Treff in Frankfurt am Main** – the prelude to everything that followed: the memoranda series and the [SituationReport software](#). It is based on the original slide deck and the transcript of the recording (archived in the curator's working archive). Part A reproduces the talk in a smoothed but substantively unchanged form; Part B looks back from today's vantage point (August 2026); **Appendix A** shows the original slides, **Appendix B** introduces each memorandum published since then on one page. This is **Version 1.3** (Appendix B updated to the AI memorandum v2.0). Write-up and contextualization with the support of **Claude** (Anthropic – Claude Fable 5, via Claude Code); editorial responsibility: Robert Seebauer. Earlier editions: [release archive](#) [denkschriften/](#)<date> [on GitHub](#).

Executive Summary

What was it about? A gap. Between the strategy level of large enterprises (comprehensive, direction-setting, goal-setting – but without ground contact) and the tactical level (which does the actual work – but optimizes locally), the layer the military has spent two centuries professionalizing is missing: **operational planning**. And its first tool is missing right along with it: an honest **situational picture**.

THE SINGLE MOST IMPORTANT SENTENCE IN THIS DOCUMENT

Between strategy and tactics, operational planning is missing – and its first tool is the situational picture: **if you do not know where you stand, you need not tell anyone where to march.**

The three core messages of the talk:

- 1 The missing layer.** Setting goals and trusting that “people will do the right thing anyway” works just as little as *Microcontrolling* by a leadership that cannot possibly possess the expertise of hundreds of thousands. The military's answer since the Prussian reforms: operational planning as a **professionalized craft** instead of hoping for the genius of a single individual.
- 2 First the situational picture, then the plan.** Leading begins not with goals but with the situation. Today's “situational pictures”, however, are slide battles (McChrystal's spaghetti slide), punchy KPIs with no message (the streetlight effect) and 20-minute discussions about €35 in a €120 million portfolio. What we would instead *at least need to know* fits into four fields: **Flow, Quality, Value, Dependencies**.
- 3 If it isn't (yet) possible: skill instead of rules.** Where no operational planning is possible, the rule is: the higher the uncertainty, the *fewer* the rules – replace rules with the skill of those doing the work, give a goal instead of instructions, and train for real (10–20% of working time, not one web-based training per year).

Why “Memorandum 0”? Because the talk ended with an announcement: “*Today's talk is really just the prelude to a series of memoranda ... alongside that, I'll be working on my software, which will then support all this.*” Both have come true – this document places the origin at the head of the series.

Part A – The talk (curated version)

A.1 Two levels – and nothing in between (slides 2–3)

Large enterprise organizations operate on a two-level pattern. At the top, **strategy**: big thoughts, big wishes, in the best case for the good of the company – all-encompassing, direction-setting, goal-setting. At the bottom, **tactics**: the discipline of actually doing things – how to deploy software, how to design testing, how to close the books. That is where the people work who actually take part in the product.

The problem: **no real transition** takes place between the two. Strategy produces castles in the air – “ideally 100 goals at once, because *one* goal would be far too goal-oriented” – and tactics, meanwhile, does what its training and experience dictate. Both fail predictably:

- **Merely setting goals** and trusting doesn't help – the teams optimize locally without noticing that they are working against other teams, or that the same thing is being built three or four times in parallel (often with two more copies commissioned externally on top).
- **Microcontrolling** helps even less: a board cannot look over the shoulders of five hundred thousand employees, and it cannot possess the expertise of a thousand specialist fields. The prince who runs after every single subordinate to correct him leads his country down the drain.

The solution is as old as the problem: everyone who ever had to coordinate large formations – Caesar, Charlemagne, every Chinese emperor – needed people to do it for him. And so that this would not depend on the **genius of a single individual**, it was professionalized: after its defeats in the Napoleonic Wars, Prussia introduced **general staff training** – the insight that, without Old Fritz, you must at least reliably organize *average to above-average* instead of hoping for the next genius. At its core, a general staff does nothing but **operational planning**: it takes up strategy's big goals, understands them, and breaks them down into achievable intermediate goals, plans and missions – not as “do this, do that”, but as goals with proposals for who needs to work with whom.

Exactly this layer is missing in corporate IT. SAFe calls it Large Solution, and many companies have formally established it – but what is actually lived there is usually a pass-through of goals, a budget check, a controlling mechanism. Real situation assessment, forging of plans, aligning, communicating cross-cutting initiatives? In reality, it does not happen.

A.2 Thesis: before any leadership comes the situational picture

Why does the solution level not work? The talk's first thesis: **forging plans and communicating goals only makes sense if you know what the situation is.** Figuratively: there is no point in telling anyone to march to the capital if I do not know where we are, where the capital lies, and how to get there. Only once the situational picture stands – *here* is where we are, *there* is where we want to go – can leadership sketch the transition and leave it to tactics *how* to march.

A.3 What situational pictures look like today (slides 4–5)

The talk showed the most famous slide in military history: the spaghetti diagram of Afghan insurgency dynamics that General **Stanley McChrystal** was handed by his staff – hundreds of nodes and arrows, absolutely nothing to make out. His recorded comment: “*When we understand that slide, we'll have won the war.*” In 2010 the slide made it into the New York Times, under the headline “**We Have Met the Enemy and He Is PowerPoint**”. The finding behind it is serious: young officers – in our world, young team leads – spend their days *crafting* tomorrow's daily situation briefing instead of *leading*. Leadership capacity is briefed away.

And in our world? Three patterns, familiar to every portfolio practitioner:

- **The streetlight effect.** The drunk looks for his key under the streetlight – not because he lost it there, but because that is where the light is. That is exactly how we measure: we delight in the scrap of information that happens to be lit and never ask whether it has any relevance to the overall problem. The monsters stay in the dark – and as long as we don't look at them, they're not there.
- **Punchy but meaningless KPIs.** By definition, a key performance indicator should indicate the company's success. Typical transformation KPIs: number of trained Scrum Masters. Lines of code. Looks good, says nothing.
- **The €35 discussion.** Twenty minutes of debate about a snack bill – in a portfolio with a €120 million budget, while decisions about hundreds of thousands were pending and got postponed by six weeks. We burn time on discussions that bear no relation to what we do *not* discuss.

A.4 What would we at least need to know? (slide 6)

The constructive core of the talk: if we wanted to build a situational picture – what would it need to answer *at a minimum*? Four fields, sixteen guiding questions:

Field	Guiding questions
Flow is the work flowing?	How long do we take for a feature (on average, with what spread)? How many features do we complete per time slice (day/week/PI)? How stable are our forecasts – do our estimates hit the mark, and do we know it? Is our WIP healthy – utilized with a buffer, or long since overloaded because we attempt everything at once?
Quality what comes out the other end?	Do we even have quality goals – do we know what “good” means for <i>us</i> (Bentley and Dacia both comply with the regulations and yet mean completely different things)? Do we meet them? Are we sustainable – will our product survive the next five years of change? What does a deliberate reduction in quality cost us – and which technologies do we really need (instead of gold-plated solutions with menu items no one has visited since Dr. Müller retired)?
Value is it worth it?	How large is the Cost of Delay (Reinertsen: if you only quantify one thing, quantify the cost of delay)? What happens to all running initiatives when we reprioritize – does the board's new darling eat up its own gains through the delay of all the others? Are our solutions even profitable (or does the €25 million cloud solution merely replace the value the old software already delivered)? Where do we need to invest – including research <i>and</i> the never-budgeted decommissioning costs?
Dependencies how does it all hang together?	How do our products and applications hang together? Which skills do we need to make all of this happen? What has to go, what may live on (what have we implemented three times)? Where are risks lurking – identified, assessed, treated, tracked – and what are we doing about them?

Throughout, the yardstick is the operational planner's: not perfection, but **workability**. “50 days on average, spread 40–60, 85% within the corridor” is a statement you can plan with – “never looked at it, maybe in two years” is not. An estimation accuracy of 50% is neither good nor bad; *not knowing it* is bad.

A.5 Who does this? (slide 7)

In SAFe, the roles already exist: **Solution Train Engineer** (flow, coaching, topologies), **Solution Management** (Business Capability Map, business runway, business vision, Horizon Planning), **Solution Architect** (Architectural Runway, skill map, technical roadmap, technical vision). The tasks from A.4 can be cleanly assigned to these three roles – it takes **no fourth role and no expensive consulting firm**.

But: three people for 14 ARTs, 2,400 employees and 40 products cannot do this alone. “*That support can either be other people – or maybe something like AI, maybe some kind of great reporting tools.*” And if the three of them can't manage it time-wise: first, stop taking care of things that are none of their business – second, reinforcements.

A.6 And what if I can't or am not allowed to? (slide 8)

What if the operational-planning layer cannot be pushed through politically, the people are missing, or the system is simply too **complex** to survey from above? Then three pieces of advice – which are really one:

1. **The higher the uncertainty, the fewer the rules.** In a crisis, companies do the opposite: more guidance, more compliance guardians, more defined processes. That is exactly the wrong way round – it is a fallacy to believe uncertainty can be foreseen and regulated away. Making decisions without a reliable situational picture and pushing them through by decree is worse than accepting your own ignorance and trusting those who do the work.
2. **Replace rules with skill.** Those who do the right thing need no instructions – they need a **goal**. That requires training on the armies' principle: everyone is trained **one level above** where they stand – so they can step in, and so they know *how their boss thinks* when they have to decide without checking back.
3. **Train for real.** “One web-based training a year is not training – it's bullshit.” What is meant is 10–20% of working time – half a day to a full day per week – of guided learning, trying things out, exchanging experience. The model is two millennia old: Flavius Josephus on the Roman army – “*Their drills are bloodless battles, and their battles bloody drills.*” For those who drill like this, the real thing is just another drill – serious this time. (With a wink, the slide in the talk attributed the quote to a “manager at Netflix (or something)”.)

A.7 The close of the talk

“Today's talk is really just the prelude to a series of memoranda I intend to write – how does solution management actually work, what would I need for it, how can I support it, how do I tie it in with today's lean-agile frameworks. Alongside that, I'll be working on my software, which will then support all this – a tool, a tool suite for it. That is what will be coming in the future.”

Part B – What became of it (retrospective, as of August 2026)

Six months later, the prelude has become a program. Every core thesis of the talk has since received its full elaboration – and several asides turned out to be seeds:


Thesis/motif of the talk	What became of it
The missing operational-planning layer ; the general staff as the model	ToT memorandum : Shared Consciousness + Empowered Execution, staff craft, 12-month curriculum. EA memorandum : the solution level as an architectural level of responsibility in the control loop
First the situational picture, then the plan	The guiding idea of the entire series – and of the SituationReport software : the situational picture as the carrier of the control loop (the EA memorandum's single most important sentence), the ritual as its engine (the ToT memorandum's)
The four knowledge fields Flow · Quality · Value · Dependencies	Flow : implemented as flow metrics (<code>build_reports</code> , <code>portfolio</code> , Monte Carlo <code>simulate</code>). Value/Dependencies : EA dimensions of roadmap Phase B (capability health, ROAM, dependency register, decision log). Quality : NFR/runway register (Phase B) + SLO/DORA (Phase C). 16 guiding questions became a software roadmap
The McChrystal slide and "Death by PowerPoint"	Became the core source: <i>Team of Teams</i> is the foundation of the ToT memorandum – the O&I briefing as the counter-model to the slide battle, the situation report as its artifact
Estimate ranges instead of point values ("40–60 days, 85% within the corridor")	The Simulate module calculates exactly that way: Monte Carlo confidence intervals instead of point forecasts; Hubbard and Duke later supplied the measurement-science and judgment-psychology foundations
"Maybe something like AI, maybe some kind of great reporting tools"	Both came true: the reporting tool suite (SituationReport) and the AI memorandum : "AI is a staff member, not a commander – from report to situation service", including software Phase D
Training instead of WBT , thinking one level up, Josephus	ToT memorandum Parts B–D: staff training as a craft, competency fields 1–7, a curriculum with katas and planning exercises – "drill-based like the Roman army" became the program
Cost of Delay (Reinertsen), Horizon Planning, decommissioning costs	Anchored in the situational-picture literature list and the EA dimensions; WSJF/Lean-budget topics of the program tables
"The higher the uncertainty, the fewer the rules"	Mission command (Auftragstaktik) as the common thread of the series (Clausewitz → Auftragstaktik → Marquet → McChrystal); defended in the AI memorandum against the recentralization reflex

The punchline at the end: the talk promised a series of memoranda and a piece of software. As of August 2026 there are four documents (ToT · EA · side document Architecture Thought Leaders · AI), a literature list with five additions, a curriculum, a software product with seven modules including a roadmap – and this document, which records the beginning.

A note on the write-up: The transcript is an audio recording with the usual mishearings – General **McChrystal** became “Jeremy Crystal”, **Reinertsen** became “Don Reineson”; both are silently corrected here. The NYT article on the spaghetti slide (“We Have Met the Enemy and He Is PowerPoint”, Elisabeth Bumiller) appeared on April 26, 2010 – dated freely from memory in the talk. The Josephus quote (*Bellum Judaicum*, c. AD 75–79) has come down to us in paraphrase; the Netflix joke comes from the original slide. All theses, examples and formulations are otherwise unchanged from the talk.

Appendix A – The Original Slides

The complete slide deck of the talk (9 slides), exported unchanged from the original. The slides are deliberately a document of their time – including the typo (“Porfolio”), the deck’s placeholder date and the Netflix joke.




Robert Seebauer

Large Solution Management

Alles wird gut, oder doch?

14.08.2026 1

Slide 1 – Title: “Large Solution Management. Everything Will Be Fine, Right?”



Warum Large Solution?

Portfolio

Strategie
Zum Wohle des Unternehmens, umfassend, wegweisend, zielgebend

Hier fehlt was

ART

Taktik
Wie arbeiten wir?, Handarbeit, tagesaktuell

14.08.2026 Robert Seebauer 2

Slide 2 – Why Large Solution? Strategy and tactics – “Something is missing here”

Warum Large Solution?

Portfolio

Strategie

Zum Wohle des Unternehmens, umfassend, wegweisend, zielgebend

Solution

Operationsplanung

Wie erreichen wir das Ziel? kombinierend, unterstützend, Schwerpunktbildung

ART

Taktik

Wie arbeiten wir?, Handarbeit, tagesaktuell

14.08.2026


Robert Seebauer

3

Slide 3 – The missing layer: operational planning at the solution level

Wie läuft's so?

Afghanistan Stability / COIN Dynamics



"When we understand that slide, we'll have won the war,"

14.08.2026

Robert Seebauer

→ We Have Met the Enemy and He Is PowerPoint

4

Slide 4 – "We Have Met the Enemy and He Is PowerPoint" (McChrystal's spaghetti slide)

Oder doch eher so?

- Fixierung auf **nichtssagende**, aber dafür markige KPI
- "Betrunkener sucht Hausschlüssel" Phänomen
- 3000 EUR Diskussionen bei 200 MEUR Budget




14.08.2026

5

Slide 5 – Or rather like this? KPI fixation, the streetlight effect, €3,000 discussions

Was müssten wir wenigstens wissen?




14.08.2026

Robert Seebauer

6

Slide 6 – What would we at least need to know? Flow · Quality · Value · Dependencies

Wer macht das?



14.08.2026

Robert Seebauer

7

Slide 7 – Who does this? STE, Solution Management, Solution Architect

Und wenn ich das nicht kann / darf?



Je höher die Unsicherheit, desto weniger Regeln

Ersetze Regeln durch Können der Ausführenden.

Wer das richtige tut, braucht keine Anweisungen.

Er braucht ein Ziel

Einmal im Jahr ein Web Based Training ist keine Ausbildung, sondern Bullshit. **Jede Woche 4h echtes Training** sind eher brauchbar.



14.08.2026

Robert Seebauer

*Flavius Josephus, Manager bei Netflix (oder so)

8

Slide 8 – And what if I can't or am not allowed to? Situational picture, skill instead of rules, real training

Wie geht's weiter?

Fragt

Wolfgang, Jan, Heiko, Sabine, Ilka, Richard & Manfred

Und ich?

Werde die nächsten Monate die einzelnen Bereiche
daraus tiefer legen , schaut ins Internet

Appendix B – The Series at a Glance: each memorandum on one page

Four documents have appeared since the talk. Here, each on one page – as an invitation to read on.

Team of Teams & the Training of Staffs and Operational Planners

Version 4.1 · 2026-08-14 · 19 pages · DE + EN

THE SINGLE MOST IMPORTANT SENTENCE

It is not the reporting tool that creates the operating picture, but the *ritual* and the *trained people* around it – the tool merely makes the ritual cheaper, faster and more honest.

What it is about

What portfolio and solution levels can learn from McChrystal's Team of Teams and 200 years of staff training: a shared operating picture plus decentralized decision authority – and planning as a teachable craft instead of an act of heroism.

What is inside

- Shared Consciousness + Empowered Execution – effective only as a pair; the O&I briefing as the model ritual
- Staff craft in plain language: MDMP, wargaming, red teaming, AAR, backbrief
- General staff training in DE/US/UK/AT – staged pipeline, selection, assignment coupling
- Five complementary core sources (Clausewitz, Marquet, Pink, Hubbard, Duke) + four civilian corroborating witnesses

Directly usable

- The 12-month “Portfolio Staff Training” curriculum · a five-step quick start for those in a hurry
- A program table of purchasable trainings · the anti-resulting rule for the next AAR

→ Download (DE + EN): [Online overview of the series](#)

Enterprise Architecture & the Situational Picture at the Solution Level

Version 2.1 · 2026-08-14 · 10 pages · DE + EN

THE SINGLE MOST IMPORTANT SENTENCE

The solution level is not a “coordination layer” but the **architectural responsibility level**: it keeps the control loop between strategy (EA) and delivery (the ARTs) closed – and the **situational picture is the instrument** with which it does so.

What it is about

What portfolio and solution leaders can learn from the discipline of enterprise architecture – and what content a solution situational picture needs beyond the flow metrics.

What is inside

- The EA ↔ solution control loop: frame downwards, reality upwards · capabilities as the shared language · the architect as the “Rosetta Stone”
- Five situational-picture principles: capability before feature · debt is a first-class citizen · two readings, one picture · confidence belongs to the number · the situational picture is a governance instrument
- Four new report dimensions: capability health · NFR/runway · ROAM risks · dependencies
- Four EA works (Graves, Carducci, Gu, Milchman) plus four thought leaders in profile

Directly usable

- 90-day entry path: capability map → ROAM/decision log → two-readings report
- Software roadmap Phases A–C (organized by data origin)

→ Download (DE + EN): [Online overview of the series](#)

Four Architecture Thought Leaders: Hohpe · Fowler · Martin · Hohmann

Version 1.1 · 2026-08-14 · 12 pages · DE + EN

THE SINGLE MOST IMPORTANT SENTENCE

Four paths, one destination: architecture work is **deciding under uncertainty with a shared language** – its tools (patterns, options, boundaries, games) are therefore **leadership and staff tools**, not technology internals.

What it is about

The detailed elaboration on the four thought leaders who carry the ToT and EA memoranda as corroborating witnesses – their works, core concepts, and what they yield for the situational picture, staff work and the software.

What is inside

- Hohpe: Architect Elevator, “Architects live in the first derivative”, Enterprise Integration Patterns, Cloud/Platform Strategy
- Fowler: the PoEAA pattern language, refactoring & the debt quadrant, Strangler Fig
- Martin: Clean Architecture, the Dependency Rule, katas & a culture of practice
- Hohmann: Marketecture, Innovation Games, Participatory Budgeting, Profit Streams

Directly usable

- A synthesis with six software derivations – among them trend columns per metric, the debt quadrant in the register, the strangler progress metric, a profit-stream field per capability

→ Download (DE + EN): [Online overview of the series](#)

AI and the Situational Picture – From Report to Situation Service

Version 2.0 · 2026-08-15 · 19 pages · DE + EN

THE SINGLE MOST IMPORTANT SENTENCE

AI makes situational information cheap while judgment remains scarce – therefore AI belongs on the staff, not on the commander's hill, and its greatest lever is **provision**: the right information, at the right time, in the language of the recipient.

What it is about

To what extent AI helps with creating a situational picture and above all with the provision of situational information – with military intelligence as the reference practice and the cyber SOC as the civilian proof of existence.

What is inside

- The four-level taxonomy (rules · statistics/ML · GenAI · agentic AI) · the intelligence cycle as the backbone · four AI roles on the staff · the automation dial
- **Seven situation-service patterns**: readings, delta briefing, queryable situational picture, wake-up call, delivery, situation briefing, memory
- **Seven AI principles** – among them “The LLM writes, it does not calculate”, “The situational picture knows teams, not individuals”, “Judge independently first, then ask the AI”
- Since v2.0 the **judgment block**: decomposing every decision into six components (Prediction Machines), bias ≠ noise and decision hygiene (Noise), the jagged frontier and the human in the loop (Co-Intelligence)
- NATO/RAND reference practice, DORA evidence, a risk/governance catalogue with a legal review box

Directly usable

- The 90-day pilot (delta briefing first, with a noise audit) · mapping software/AI/human per decision component (C.6) · competency field 7 and curriculum module 7 · software Phase D

→ Download (DE + EN): [Online overview of the series](#)

Sources

- **Original material:** slide deck [Robert Seebauer - Solution Management Was brauchen wir dazu.pptx](#) (9 slides, reproduced in full in Appendix A) and transcript [Vortrag Februar 2026.docx](#) – archived in the curator's working archive; talk delivered on February 26, 2026, at the SAFe Treff in Frankfurt am Main.
- Bumiller, Elisabeth: *"We Have Met the Enemy and He Is PowerPoint"*. *The New York Times*, April 26, 2010 – the McChrystal spaghetti slide.
- Reinertsen, Donald G.: *The Principles of Product Development Flow* (2009) – Cost of Delay.
- Flavius Josephus: *De Bello Judaico* (Book III) – "bloodless battles"; [free English translation \(Project Gutenberg\)](#).
- **The series that grew out of it:** Team of Teams & Staff Training · Enterprise Architecture & the Solution Situational Picture · Architecture Thought Leaders · AI and the Situational Picture · the situational-picture literature list · software roadmap [docs/portfolio_roadmap_solution-lagebild.md](#).

Version history

Version	Date	Changes
1.0	2026-08-13	First version: curated write-up of the February 2026 talk (Part A, along slides 2–9) + retrospective "What became of it" (Part B, as of August 2026) + notes on the transcript corrections; placed at the head of the series as Memorandum 0.
1.1	2026-08-13	Additions on request: event context (SAFe Treff, Feb 26, 2026, Frankfurt am Main); Appendix A with all 9 original slides; Appendix B "The Series at a Glance" – one flyer page per published memorandum (ToT v4.0, EA v2.0, companion memorandum v1.0, AI v1.1).
1.2	2026-08-14	Reference maintenance for the online publication: series cross-references now link the online overview, freely available sources link their real URLs, workshop paths removed; cross-references now without version numbers. Content unchanged.
1.3	2026-08-15	Appendix B: flyer of the AI memorandum updated to v2.0 (seven AI principles, judgment block, C.6, noise audit). Otherwise unchanged.

Memorandum 0 "Large Solution Management – Everything Will Be Fine, Right?" · v1.3 · 2026-08-15 · Talk: Robert Seebauer, SAFe Treff Frankfurt am Main, Feb 26, 2026 · Write-up with Claude Fable 5 (Anthropic) · Current editions online: <https://jaegerfeld.github.io/situation-report/denkschriften/>