Asking the Situational Picture (D3)

14 Oct 2026 · One question about one stand – every answer cites the data path, or it is not written.

What is it?

Whoever reads a situational picture has questions: which dependency is overdue? Where does Alpha-2 stand on the median? Until now the answer was: search. Now the stand can be asked – and the AI may only say what it can cite.

What it does

• The stand becomes a fact index: every number and register entry gets a path such as S1/total.median_ct or S1/governance.dependencies(D-7).status
• The index is the model's only context – no vector store, no web search, no external sources
• Citation guard: every statement cites a path, every number comes from the cited facts – otherwise the answer is discarded
• 'The situational picture holds no information on that.' is the permitted answer when the facts are silent
• Own assessment first: the field comes before the answer, never reaches the model and sits next to the answer in the file

Three decisions

Facts, not full text. The index is small and deterministic; embeddings, chunking and the limits of RAG fall away. The German DSK guidance (10/2025) recommends exactly the principle: knowledge from the reference data, the model only for language.
Hard, not soft. An uncited statement is not flagged but discarded – the same line as the numbers guard and llm check.
Explain, do not judge. D3 belongs to prediction/feedback. No recommendation, no judgement, no chat memory – every question anew against the stand, so that D7 can later measure the spread.

How to use it

```
python -m portfolio solution.json --ask "Which dependency is overdue?" --narrate ollama --llm-lang en
python -m portfolio solution.json --ask "..." --own-view "I suspect D-7." --since-snapshot prev.json
python -m portfolio solution.json --ask "..." --narrate export   # prompt bundle; return path: llm check
```