Data Minimisation at the Source (R1)

15 Oct 2026 · The picture knows teams, not persons – now already in the raw file, not only in the report.

What is it?

No module of the pipeline reads a person field. Yet every Jira export carried authors with name and e-mail in every changelog entry, and the manual export added assignee, reporter and comments. Rohrlich (“KI und Recht”) extends our rule from the output to the input: minimise data as early as possible.

What it does

• Allowlist: only the seven fields transform_data reads; --keep for further fields without personal reference
• The changelog stays complete – every history, every item; authors become a constant placeholder, person items lose their values, free-text items their texts
• The REST fetch always minimises; manual exports via python -m get_data minimise, check warns, GUI button Minimise
• A report of what was removed; contract test: raw and minimised give the same figures

Three decisions

Allowlist, not denylist. A list of known person fields misses custom fields and free text with names. What is not needed falls.
Anonymise, do not pseudonymise. A per-person pseudonym would still be personal data. The placeholder is the same for everyone – the transitions stay gapless all the same.
No opt-out on fetch. No purpose of the pipeline needs persons. This does not replace the legal basis – the docs say so as an assessment, not as legal advice.

How to use it

```
python -m get_data fetch --url https://company.atlassian.net --project ART_A --email name@company.com --
python -m get_data minimise ART_A_merged.json   # -> ART_A_merged_min.json
```