

LLM Executive Summary

Feature D1 · The language model phrases the report's management summary — numbers come only from the data path

What is it?

The report shows the management summary as a table; the executive summary turns it into 6–9 sentences of management prose: overall picture, standout units, data confidence, governance head counts. Phrased by the exchangeable llm framework (ollama local, claude external, mock for demos) — the input is exclusively the deterministic metrics contract produced by the same data path as report and snapshots.

How to use it

```
python -m portfolio my_solution.json --output report.html --narrate
# section 'Executive Summary (Entwurf)' below the summary table
# + report.html.exec_summary.md for editing
# + llm_audit.jsonl (purpose d1_exec_summary)
```

In the GUI, the existing 'AI narration (draft)' checkbox now also covers 'Generate report ...'. If the AI fails, the report is written without the draft — clean degradation.

The guards (unchanged)

- Numbers guard: every number verbatim in the contract, otherwise the text is discarded
- Art. 50 banner: model, deployment class, prompt version — never claims approval
- Operator evidence: llm_audit.jsonl with SHA-256 hashes, never full texts

Guardrails

The contract names units and register counts — owner/person fields structurally never appear (enforced by tests). Without --narrate the report stays byte-identical; PDF runs get a hint instead of a draft.