llm check – the return path of the prompt export

14 Oct 2026 · A foreign answer enters the report only after passing the same guards as a draft produced inside the tool.

What is it?

Since 0.42.0 every AI step can export its prompt to be run on another machine or model. The answer came back by hand – unchecked. The export was half a road. llm check is the way back.

What it does

• Reads the answer from a file or the clipboard and finds the matching bundle (newest prompt_*.md next to it, or --bundle)
• Numbers guard against the contract in the bundle; questions guard for red-team answers
• Pass: draft with the banner 'Externally produced AI draft' as .checked.md; reject: nothing written, exit code 2
• Audit record with the same input hash as the export – outward and return path are linked in the evidence

Three decisions

Reject, do not flag. An invented number does not get through, not even with a warning. The same line as a direct model call.
An honest deployment class. The tool never saw the model, so the class is external_answer and the model name comes from the operator – labelling follows what is known (D9).
Checksum in the bundle. The contract sits verbatim in the bundle, its hash in the head. A bundle edited after export is reported.

How to use it

```
python -m portfolio solution.json --narrate export --output report.html
# ... run the prompt elsewhere, save the answer as answer.md ...
python -m llm check answer.md --model "gpt-something"
python -m llm check --clipboard --bundle prompt_d2_narration_20261014-0900.md
```