

Choosing a Model and an Address

More than one model in Ollama – now selectable in the GUI too

What is it?

The GUI let you pick the PROVIDER (ollama / claude / mock) but not the MODEL. With several models pulled in Ollama, the GUI always ran the default – the `--llm-model` flag existed only on the command line. And the service address lived in the code alone: Ollama on another machine was unreachable without editing source.

What works now

- A Model field next to the provider picker: offers the installed models but stays typable – a model still being pulled can be entered already
- An Address field: Ollama on another machine. Without one, the `OLLAMA_HOST` environment variable applies, else `http://localhost:11434`
- A new command, `python -m llm models`: lists what is actually installed and marks the default
- `--base-url` and `--llm-base-url` on the command line

Two design decisions

These settings belong to the machine, not to the portfolio. A solution configuration is saved with relative paths and handed to colleagues – an Ollama address from someone else's laptop would be a puzzle inside it, not a setting. Provider, model and address therefore go into the machine's preferences and come back on the next start.

Nothing is silently substituted. A model you typed stays, even when it is not in the list – it may simply not be pulled yet. Ollama then says which ollama pull command is missing, which is better information than a GUI quietly running something else. If Ollama is not running, the field stays an ordinary text box – no wait, no error dialog.

How to use it

```
python -m llm models
python -m llm test --model qwen3.8:latest
python -m llm models --base-url http://gpu-box:11434
```