

AI Narration

Feature D2 part 2 · A language model drafts the situational narrative — local first, always labeled

What is it?

The delta briefing delivers the facts; the narration turns them into a 5–8 sentence draft for the Value-Stream Conference: what changed, what deserves attention — worsenings first. Phrased by an exchangeable language model, local by default (Ollama, mistral-nemo): data never leaves the machine, no account, no API approval. Alternatives: claude (Anthropic API) or mock (a stand-in without any model, for demos and tests).

How to use it

```
python -m portfolio --delta before.json today.json --narrate --output delta.html
python -m llm providers      # inventory
python -m llm test           # wiring check after installing Ollama
```

Or via the GUI: tick 'AI narration (draft)' and pick the provider next to 'Delta briefing ...'; the test-data generator shows the complete flow through the mock provider — no installation at all. Setting up Ollama: separate guide [ollama_Installationsanleitung_EN.pdf](#) (docs site → Tutorials).

Three guards, wired in

- Numbers guard: every number must occur verbatim in the briefing, otherwise the text is discarded — 'the LLM writes, it does not calculate'
- AI label (Art. 50 EU AI Act): every draft visibly carries model, deployment class and prompt version — only the editing human can approve
- Operator evidence: `llm_audit.jsonl` logs every request with SHA-256 hashes — never full texts, never keys

Guardrails

Without `--narrate` the briefing stays exactly as before — AI is an add-on, never a prerequisite. The model sees only the delta Markdown (teams, never persons); a provider is ONE file in `llm/providers/` — exchangeable like the data sources. API keys only from environment variables.