Noise Audit (D7)

14 Oct 2026 · The same question five times – how much does the model scatter before you trust its answer?

What is it?

A language model does not answer the same question the same way twice. Taking one answer as objective imports that spread as false objectivity. The noise audit makes the model's loudness measurable – on your own situational picture.

What it does

• Asks the question n times in varied phrasings against the same fact index (run 1 unchanged, then your variants, then fixed frames)
• Measures from the citations: consistency 0–100 (Jaccard), core (majority), scatter (one answer only), number conflicts per path
• Counts 'no information' and rejected answers as rates – rejected is not retried
• Lists staff assessments next to it, unscored

Three decisions

No second model. Similarity is computed from the cited paths, not estimated by an AI – otherwise you measure its spread too.
Rejected counts. When a rephrasing makes the model fall out of the contract, that is a finding, not an outlier to retry away.
Finding, not judgement. The protocol never says 'too high'. What follows from the spread – switch model, sharpen the question, stay deterministic – is a human decision.

How to use it

```
python -m portfolio solution.json --ask "Which dependency is overdue?" --noise 5 --narrate ollama --llm-
python -m portfolio solution.json --ask "..." --noise 7 --noise-variants variants.txt --views staff.txt
```