

Red-Team Assistant

Feature D5 · Premortem questions from the decision log — raw material only, never judgements

What is it?

Good staffs attack their own decisions and assumptions before reality does — but red-team capacity is scarce. D5 turns the decision/assumption log (B4) into premortem and attack questions as raw material for a humanly moderated session: for decisions the look back from the future ('what made this fail six months from now?'), for assumptions the direct attack ('what would have to be true for it to flip — and how would one notice early?').

How to use it

```
python -m portfolio my_solution.json --red-team questions.md
# 1-3 questions per log entry, grouped by ID,
# with the Art. 50 banner + llm_audit.jsonl (d5_red_team)
```

Provider as everywhere: local ollama (default), external claude, mock for the model-free demo. Prerequisite is a referenced decision log (config field decisions) — if it is missing, the error says exactly that.

The questions guard

The KI-Denkschrift maps D5 to JUDGEMENT — and demands: raw material only, no recommendation button. That is machine-enforced here: every '-' line of the output must end as a question, otherwise the whole draft is discarded — on top of the numbers guard, labeling and audit. The tool CANNOT deliver recommendations.

Guardrails

Questions name teams and entries, never persons; numbers and IDs come verbatim from the log. Moderation, selection and answering stay entirely with the humans.